

Operating Systems Principles

UCLA-CS111-W18

Quentin Truong
Taught by Professor Reiher

Winter 2018

Contents

| | | |
|----------|--|----------|
| 1 | L3: Arpaci-Dusseau Chapter 5: Interlude: Process API | 2 |
| 1.1 | The fork() System Call | 2 |
| 1.2 | The wait() System Call | 2 |
| 1.3 | The exec() System Call | 2 |
| 1.4 | Why? Motivating The API | 2 |
| 1.5 | Other Parts Of The API | 2 |
| 2 | L3: Arpaci-Dusseau Chapter 6: Mechanism: Limited Direct Execution | 2 |
| 2.1 | Basic Technique: Limited Direct Execution | 2 |
| 2.2 | Problem 1: Restricted Operations | 3 |
| 2.3 | Problem 2: Switching Between Processes | 3 |
| 2.4 | Worried About Concurrency? | 4 |
| 2.5 | Summary | 4 |
| 3 | L3: Linking and Libraries: Object Modules, Linkage Editing, Libraries | 4 |
| 3.1 | Introduction | 4 |
| 3.2 | The Software Generation Tool Chain | 4 |
| 3.3 | Object Modules | 4 |
| 3.4 | Libraries | 5 |
| 3.5 | Linkage Editing | 5 |
| 3.6 | Load Modules | 5 |
| 3.7 | Static vs. Shared Libraries | 5 |
| 3.8 | Dynamically Loaded Libraries | 5 |
| 4 | L3: Linkage Conventions: Stack Frames and Linkage Conventions | 5 |
| 4.1 | Introduction | 5 |
| 4.2 | The Stack Model of Programming Languages | 5 |
| 4.3 | Subroutine Linkage Conventions | 6 |
| 4.4 | Traps and Interrupts | 6 |
| 5 | L4: Arpaci-Dusseau Chapter 7: Scheduling: Introduction | 7 |
| 5.1 | Workload Assumptions | 7 |
| 5.2 | Scheduling Metrics | 7 |
| 5.3 | First In, First Out (FIFO) | 7 |
| 5.4 | Shortest Job First (SJF) | 7 |
| 5.5 | Shortest Time-to-Completion First (STCF) | 7 |
| 5.6 | A New Metric: Response Time | 7 |
| 5.7 | Round Robin | 7 |
| 5.8 | Incorporating I/O | 8 |

| | | |
|-----------|---|-----------|
| 5.9 | No More Oracle/Summary | 8 |
| 6 | L4: Arpaci-Dusseau Chapter 8: Scheduling: The Multi-Level Feedback Queue | 8 |
| 6.1 | MLFQ: Basic Rules | 8 |
| 6.2 | Attempt 1: How To Change Priority | 8 |
| 6.3 | Attempt 2: The Priority Boost | 8 |
| 6.4 | Attempt 3: Better Accounting | 8 |
| 6.5 | Tuning MLFQ And Other Issues | 8 |
| 6.6 | MLFQ: Summary | 9 |
| 7 | L4: Real Time Scheduling | 9 |
| 7.1 | What are Real-Time Systems | 9 |
| 7.2 | Real-Time Scheduling Algorithms | 9 |
| 7.3 | Real-Time and Linux | 9 |
| 8 | L5: Arpaci-Dusseau Chapter 12: A Dialogue on Memory Virtualization | 10 |
| 8.1 | Overview | 10 |
| 9 | L5: Arpaci-Dusseau Chapter 13: The Abstraction: Address Spaces | 10 |
| 9.1 | Early Systems | 10 |
| 9.2 | Multiprogramming and Time Sharing | 10 |
| 9.3 | The Address Space | 10 |
| 9.4 | Goals | 10 |
| 10 | L5: Arpaci-Dusseau Chapter 14: Interlude: Memory API | 10 |
| 10.1 | Types of Memory | 10 |
| 10.2 | The malloc()/free() Call | 10 |
| 10.3 | Common Errors | 11 |
| 10.4 | Underlying OS Support | 11 |
| 11 | L5: Arpaci-Dusseau Chapter 17: Free-Space Management | 11 |
| 11.1 | Assumptions | 11 |
| 11.2 | Low-level Mechanisms | 11 |
| 11.3 | Basic Strategies | 11 |
| 11.4 | Other Approaches | 12 |
| 12 | L5: Garbage Collection and Defragmentation | 12 |
| 12.1 | Garbage Collection | 12 |
| 12.2 | Defragmentation | 12 |
| 13 | L6: Arpaci-Dusseau Chapter 18: Paging: Introduction | 13 |
| 13.1 | A Simple Example And Overview | 13 |
| 13.2 | Where Are Page Tables Stored? | 13 |
| 13.3 | Whats Actually In The Page Table? | 13 |
| 13.4 | Paging: Also Too Slow | 13 |
| 14 | L6: Arpaci-Dusseau Chapter 19: Paging: Faster Translations (TLBs) | 13 |
| 14.1 | TLB Basic Algorithm | 13 |
| 14.2 | Example: Accessing An Array | 14 |
| 14.3 | Who Handles The TLB Miss? | 14 |
| 14.4 | TLB Contents: Whats In There? | 14 |
| 14.5 | TLB Issue: Context Switches | 14 |
| 14.6 | Issue: Replacement Policy | 14 |

| | |
|---|-----------|
| 15 L6: Arpaci-Dusseau Chapter 21: Beyond Physical Memory: Mechanisms | 14 |
| 15.1 Swap Space | 14 |
| 15.2 The Present Bit | 15 |
| 15.3 The Page Fault | 15 |
| 15.4 What If Memory Is Full? | 15 |
| 15.5 Page Fault Control Flow | 15 |
| 15.6 When Replacements Really Occur | 15 |
| 16 L6: Arpaci-Dusseau Chapter 22: Beyond Physical Memory: Policies | 15 |
| 16.1 Cache Management | 15 |
| 16.2 The Optimal Replacement Policy | 16 |
| 16.3 Replacement Policies | 16 |
| 16.4 Implementing LRU | 16 |
| 16.5 Considering Dirty Pages | 16 |
| 16.6 Other VM Policies | 16 |
| 16.7 Thrashing | 16 |
| 17 L6: Working Sets | 17 |
| 17.1 LRU is not enough | 17 |
| 17.2 The concept of a Working Set | 17 |
| 17.3 Implementing Working Set replacement | 17 |
| 17.4 Dynamic Equilibrium to the rescue | 17 |
| 18 L7: Arpaci-Dusseau Chapter 25: A Dialogue on Concurrency | 18 |
| 18.1 Dialogue | 18 |
| 19 L7: Arpaci-Dusseau Chapter 26: Concurrency: An Introduction | 18 |
| 19.1 Introduction | 18 |
| 19.2 Why Use Threads? | 18 |
| 19.3 An Example: Thread Creation | 18 |
| 19.4 The Heart Of The Problem: Uncontrolled Scheduling | 18 |
| 19.5 The Wish For Atomicity | 18 |
| 20 L7: Arpaci-Dusseau Chapter 27: Interlude: Thread API | 19 |
| 20.1 Threads | 19 |
| 21 L7: User-Mode Thread Implementation | 19 |
| 21.1 Introduction | 19 |
| 21.2 User/Kernel | 19 |
| 21.3 User/Kernel | 19 |
| 22 L7: Inter-Process Communication | 19 |
| 22.1 Introduction | 19 |
| 22.2 Simple Uni-Directional Byte Streams | 20 |
| 22.3 Named Pipes and Mailboxes | 20 |
| 22.4 General Network Connections | 20 |
| 22.5 Shared Memory | 20 |
| 22.6 Network Connections and Out-of-Band Signals | 20 |
| 23 L7: Named pipes, Send, Recv, Mmap | 21 |
| 23.1 Named Pipes | 21 |
| 23.2 Send | 21 |
| 23.3 Recv | 21 |
| 23.4 Mmap | 21 |

| | |
|--|-----------|
| 24 L8: Arpaci-Dussseau Chapter 28: Locks | 21 |
| 24.1 Locks: The Basic Idea | 21 |
| 24.2 Pthread Locks | 21 |
| 24.3 Evaluating Locks | 21 |
| 24.4 Controlling Interrupts | 21 |
| 24.5 A Failed Attempt: Just Using Loads/Stores | 22 |
| 24.6 Building Working Spin Locks with Test-And-Set | 22 |
| 24.7 Compare-And-Swap | 22 |
| 24.8 Load-Linked and Store-Conditional | 22 |
| 24.9 Fetch-And-Add | 22 |
| 24.10A Simple Approach: Just Yield, Baby | 22 |
| 24.11Using Queues: Sleeping Instead Of Spinning | 23 |
| 24.12Two-Phase Locks | 23 |
| 25 L6: Inter-Process Communication | 23 |
| 25.1 Introduction | 23 |

1 L3: Arpaci-Dusseau Chapter 5: Interlude: Process API

1.1 The fork() System Call

- Crux: How to create and control processes
- fork()
 - Creates new process; returns child's PID to parent; returns 0 to child;
 - Each has own PC, registers, address space
- Nondeterministic Behavior
 - Scheduler will decide which process to run
 - May lead to problems in multi-threaded programs

1.2 The wait() System Call

- wait()
 - Parent calls wait() to wait for child to finish execution

1.3 The exec() System Call

- exec()
 - Loads code, overwrites code segment, and reinitializes memory space
 - Takes executable name and arguments
 - Does not create a new process; transform current process

1.4 Why? Motivating The API

- Separation
 - Separating fork() and exec() allows code to alter the environment of the about-to-run program
- Example
 - Shell forks a process, execs the program, and waits until finished
 - The separation allows for things such as output to be redirected (closes stdout and opens file)

1.5 Other Parts Of The API

- kill()
 - System call sends signal to process to sleep, die, etc

2 L3: Arpaci-Dusseau Chapter 6: Mechanism: Limited Direct Execution

2.1 Basic Technique: Limited Direct Execution

- Crux: How to efficiently virtualize CPU with control
- Limited Direct Execution
 - OS will create entry for process list, allocate memory for program, load program into memory, setup stack with argc/v, clear registers, execute call to main()
 - Program will run main(), execute return
 - OS will free memory, remove from process list
- LDE good bc fast, but
 - Problem of keeping control
 - Problem of time sharing still

2.2 Problem 1: Restricted Operations

- User mode vs. Kernel mode
 - Restricted mode which needs to ask kernel to perform system calls
 - Calls like open() are actually procedure calls with trap to enter kernel and raise privilege
 - Return-from-trap is used to enter user mode from kernel and drop privilege
 - Push counters, flags, registers onto per-process kernel stack when trapping
- Trap table is used to control what code is executed when trapping
 - Trap handler used by hardware to cause interrupts
 - Telling hardware where trap table is is privileged
 - Trap handler actually uses system-call number, rather than specifying an address (another layer of protection)
- Two phases of LDE
 - At boot, kernel initializes trap table and remembers where it is

| OS @ boot (kernel mode) | Hardware |
|------------------------------------|--|
| initialize trap table | remember addresses of... syscall handler timer handler |
| start interrupt timer | start timer interrupt CPU in X ms |

2.3 Problem 2: Switching Between Processes

- How can OS regain control?
 - Because process is running, so OS is not running
- Cooperative Approach
 - System calls include explicit yield system call, transferring control back to OS
- Noncooperative Approach
 - Reboot, Timer Interrupt
- Saving and Restoring Context
 - Scheduler will choose when to switch processes

| OS @ run (kernel mode) | Hardware | Program (user mode) |
|--|---|--------------------------------|
| | | Process A |
| | | ... |
| | timer interrupt save regs(A) to k-stack(A) move to kernel mode jump to trap handler | |
| Handle the trap Call <code>switch()</code> routine save regs(A) to <code>proc-struct(A)</code> restore regs(B) from <code>proc-struct(B)</code> switch to k-stack(B) return-from-trap (into B) | | |
| | restore regs(B) from k-stack(B) move to user mode jump to B's PC | |
| | | Process B |
| | | ... |

2.4 Worried About Concurrency?

- Interrupt during interrupt?
 - Many complex things to do
 - Could disable interrupts (but this might lose interrupts), or locking schemes, etc

2.5 Summary

- Reboot
 - Good technique because restores system to well-tested state
 - OS will 'baby-proof' by only allowing processes to run in restricted mode and with interrupt handlers

3 L3: Linking and Libraries: Object Modules, Linkage Editing, Libraries

3.1 Introduction

- Process as fundamental; as executing instance of program
 - Program as one or more files (these are not the executables though)
 - Source must be translated

3.2 The Software Generation Tool Chain

- Source module
 - Editable text in some language like C
- Relocatable object module
 - Sets of compiled instructions; incomplete programs
- Library
 - Collection of object modules
- Load module
 - Complete programs ready to be loaded into memory
- Compiler
 - Parse source modules; usually generates assembly, may generate pseudo-machine
- Assembler
 - Object module with mostly machine code
 - Memory addresses of functions, variables may not be filled in
- Linkage Editor
 - Find all required object modules and resolve all references
- Program Loader
 - Examines load module, creates virtual space, reads instructions, initializes data values
 - Find and map additional shared libraries

3.3 Object Modules

- Code in multiple files
 - Because more understandable if splitting functionality
 - Many functions are reused, so use external libraries
- Relocatable object modules are program fragments
 - Incomplete because make references to code in other modules
 - Even the references to other code are only relative
- ELF format
 - Header section with types, sizes, and location of other sections
 - Code and data section to be loaded contiguously
 - Symbol table of external symbols
 - Relocation entries describing location of field, width/type of field, symbol table entry

3.4 Libraries

- Reusable, standard functions in libraries
 - Libraries not always orthogonal and independent
- Build program by combining object modules and resolving external references

3.5 Linkage Editing

- Resolution
 - Search libraries to find object modules to resolve external references
- Loading
 - Lay text and data in single virtual address space
- Relocation
 - Ensure references correctly reflect chosen address

3.6 Load Modules

- Load module requires no relocation and is complete
- When loading new module
 - Determine required text and data sizes and locations, allocate segments, read contents, create a stack segment with pointer
- Load module has symbol table to help determine where exceptions occurred

3.7 Static vs. Shared Libraries

- Static Linking
 - Many copies, so inefficient; also, permanent copy, so don't receive updates
- Shared Libraries
 - Implementations vary, but one way
 - Reserve address for libraries, linkage edit, map with redirection table, etc, more mapping
 - Efficient, but doesn't work for static data because one copy
 - But can be slow to load many libraries, and must know library name at loadtime

3.8 Dynamically Loaded Libraries

- DLL loaded once needed
 - Choose and load library, binds, use library, unload
 - Resource efficient because can unload
- Implicitly Loaded Dynamically Loadable Libraries
 - Another implementation of DLL with different pros/cons

4 L3: Linkage Conventions: Stack Frames and Linkage Conventions

4.1 Introduction

- What is the state of computation and how can it be saved?
- What is the mechanism of requesting and receiving services?

4.2 The Stack Model of Programming Languages

- Procedure-local variables
 - Stored on a LIFO stack
 - New call frames pushed onto stack when procedure called; old frames popped when procedure returns
 - Long-lived resources on heap, not stack

4.3 Subroutine Linkage Conventions

- X86 Subroutine Linkage
 - Pass parameters to be called by routine
 - Save return address and transfer control to entry
 - Save content of non-volatile registers
 - Allocate space for local variables
- X86 Return Process
 - Return value to where routine expects it
 - Pop local storage
 - Restore registers
 - Subroutine transfer control to return address
- Responsibilities split between caller and callee
- Saving and restoring state of procedure is mostly a matter of stack frame and registers

4.4 Traps and Interrupts

- Procedure call vs Trap/Interrupt
 - Procedure requested by running software and expects result; linkage conventions under software control
 - After trap/interrupt, should restore state
- How
 - Number associated with every interrupt/exception, maps to PS/PC
 - Push new program counter and program status (from interrupt/trap vector table) onto CPU stack
 - Resume execution at new PC
 - First level handler
 - Save general registers on stack
 - Choose second level handler based on info from interrupt/trap
 - Second level handler (procedure call)
 - Deal with interrupt/exception
 - Return to first level handler
 - Restore saved registers and return-from-interrupt/trap
 - CPU reloads PC/PS and resumes execution
- Stacking/unstacking interrupt/trap is 100x+ slower than procedure call

5 L4: Arpaci-Dusseau Chapter 7: Scheduling: Introduction

5.1 Workload Assumptions

- Workload as the processes running in the system
- Fully-operational scheduling discipline
 - Assume each job runs for same amount of time, arrives at same time, once started will run to completion, only uses CPU, run-time length is known

5.2 Scheduling Metrics

- Scheduling metric is something we can measure is useful for scheduling
 - $Turnaround_{time} : Time_{completion} - Time_{arrival}$
- Performance and Fairness often at odds with each other
 - Fairness measured by Jain's Fairness Index

5.3 First In, First Out (FIFO)

- Properties of FIFO
 - Simple and easy to implement while working well based on assumptions
- Convoy Effect
 - FIFO fails if few high-resource consumers are ahead of low-resource consumers

5.4 Shortest Job First (SJF)

- SJF is optimal given the assumptions
 - But fails if relaxes arrival-time assumption
 - A long process may start, then a short process comes in

5.5 Shortest Time-to-Completion First (STCF)

- Preemptive schedulers will context switch to run another process
 - Non-preemptive schedulers run jobs to completion before considering another
 - SJF is nonpreemptive
- Shortest time-to-completion (STCF) also known as Preemptive shortest job first (PSJF)
 - Anytime a new job arrives, determine which job has shortest time remaining, and runs that one

5.6 A New Metric: Response Time

- $T_{response} : T_{firstrun} - T_{arrival}$
- STCF is especially bad for optimizing response time

5.7 Round Robin

- RR (time-slicing) runs job for a time slice (scheduling quantum) before switching to next
 - Length of time slice is essential; if long, then long $T_{response}$; if short, context switching dominates
 - Must choose a length of time which will amortize the cost well
 - Also must consider cost of flushing CPU caches, TLBs, branch predictors, chip hardware
- Performs extremely poorly wrt turnaround time
 - Most fair policies (evenly distribute) are like this

5.8 Incorporating I/O

- Overlap leads to higher utilization and better performance
 - Use for IO, messages, etc
- Overlap CPU when one process requires IO
 - While IO for process A, run process B on CPU (because A is blocked)

5.9 No More Oracle/Summary

- Assumption of known run-time length is highly invalid
- Shortest job remaining optimizes turnaround time
- Alternating between jobs optimizes response time
- Looking ahead
 - Multi-level feedback: Using past events to predict future

6 L4: Arpaci-Dusseau Chapter 8: Scheduling: The Multi-Level Feedback Queue

6.1 MLFQ: Basic Rules

- MFLQ has a number of distinct queues with different priority levels
- If $\text{priority}(A) < \text{priority}(B)$, A runs
- If $\text{priority}(A) == \text{priority}(B)$, A and B run in RR
- Vary priority based on observed behavior

6.2 Attempt 1: How To Change Priority

- When job enters, has highest priority
- If job uses entire time slice, priority is reduced
- If job gives up CPU early, priority remains the same
- Assume jobs are short so that it will either complete or move down in priority
- Starvation
 - If there are too many interactive (IO) jobs, then longer processes with low priority will never run
- Gaming the scheduler
 - Could write program to use less than entire timeslice, to always keep highest priority
- Changing Behavior
 - Program may become interactive after computations, so needs higher priority

6.3 Attempt 2: The Priority Boost

- Boost all processes to top priority after a certain time length
- Difficult to know correct value for these voo-doo constant parameters (refer to Ousterhouts Law)

6.4 Attempt 3: Better Accounting

- Account CPU time (Anti-gaming method)
 - Once job uses up time allotment on given level, priority is reduced

6.5 Tuning MLFQ And Other Issues

- Difficult to find correct parameters
 - High-priority queue contains interactive processes and run for short timeslices (20ms)
 - Low-priority queue contains long-running processes and so run for longer timeslices (up to a few hundred ms)
 - Many queues, like 60

- Priorities boosted every second or so
- Other schedulers use mathematical formulas to calculate priority (decay-usage)
- Even may offer advice to scheduler using Linux's nice program

6.6 MLFQ: Summary

- Multiple levels of queues with feedback to determine priority
- Rules
 - If $\text{priority}(A) > \text{priority}(B)$, A runs
 - If $\text{priority}(A) = \text{priority}(B)$, A and B run in RR
 - When a job enters the system, has highest priority
 - When a job uses entire time allotment at a given level, its priority is reduced
 - After some time period S, move all the jobs in the system to the topmost queue

7 L4: Real Time Scheduling

7.1 What are Real-Time Systems

- Priority scheduling is best effort
 - Sometimes need more than just best effort (space shuttle reentry, data, assembly line, media players)
- Traditional vs Real-time systems
 - Turn-around time, fairness, response time for traditional
 - Timeliness may be ms/day of accumulated tardiness
 - Predictability is deviation in delivered timeliness
 - Feasibility is whether possible to meet requirements
 - Hard real-time is a requirement to run specify tasks at specified intervals
 - Soft real-time requires good response time, at the cost of degraded performance or recoverable failure
- Real-time systems
 - May know length of jobs/priorities, and starvation of certain jobs may be acceptable

7.2 Real-Time Scheduling Algorithms

- Static scheduling
 - May be possible to define fixed schedule if know all tasks to run and expected completion time
- Dynamic Scheduling for changing workloads
 - Questions of how to choose next task and how to deal with overload
- If high enough frequency of work, may just work for sufficiently-light loaded systems

7.3 Real-Time and Linux

- Linux was not designed as embedded or real-time system
 - Supports a real-time scheduler `sched_setscheduler`, but still does not have same level of response-times
- Windows believes in general throughput not deadlines, and is bad for critical real-time operations

8 L5: Arpaci-Dusseau Chapter 12: A Dialogue on Memory Virtualization

8.1 Overview

- Every address generated by a user program is a virtual address
 - Large contiguous address space is easier to work with than small crowded space
 - Isolation and protection are also important in preventing processes each other's memory

9 L5: Arpaci-Dusseau Chapter 13: The Abstraction: Address Spaces

9.1 Early Systems

- OS as set of routines (a library)
- Program in physical memory used rest of space

9.2 Multiprogramming and Time Sharing

- Multiprogramming
 - Multiple processes ready to run at a given time with OS switching between them
 - Increases utilization of CPU; increased efficiency of CPU is very relevant bc so expensive
- Timesharing and interactivity
 - Long program-debug cycles bad for programmers
 - Giving all programs full access to memory is not safe

9.3 The Address Space

- Address space is easy to use abstraction of physical memory
 - Contains code, stack, heap
 - Every program thinks it had very large address space, even though it doesn't

9.4 Goals

- Transparency
 - Cannot tell that memory is virtual
- Efficiency
 - OS should make virtualization efficient wrt time and space, relying on hardware for this
- Protection
 - Isolate process memory from each other

10 L5: Arpaci-Dusseau Chapter 14: Interlude: Memory API

10.1 Types of Memory

- Stack
 - Automatic memory is managed implicitly by compiler
- Heap
 - Long lived memory where allocations and deallocations handled by programmer

10.2 The malloc()/free() Call

- `double *d = (double *) malloc(sizeof(double));`
- `free(d);` // prevents memory leaks

10.3 Common Errors

- Modern languages have automatic memory-management or a garbage collector because people don't free
- Seg fault if you forget to allocate
- Buffer overflow if not enough allocated space
- Dangling pointer if you free memory before finished using it
- Double freeing memory is undefined
- Incorrect use of free (passing it things other than pointer from malloc) is dangerous
- Use Valgrind and Purify to find memory leaks

10.4 Underlying OS Support

- Break is the location at the end of the heap
 - System call brk is used to increase/decrease size of heap

11 L5: Arpaci-Dusseau Chapter 17: Free-Space Management

11.1 Assumptions

- Free list manages the heap; contains references to all the free chunks in the region
- External fragmentation
 - Have enough space, but not contiguous, so can't malloc
- Internal fragmentation
 - Gives memory larger than requested, which remains unused

11.2 Low-level Mechanisms

- Splitting and Coalescing
 - Split free chunk in two, returning first to the caller
 - Coalesces adjacent free memory together, forming a single larger free chunk
- Header of allocated memory
 - Contains size of region and magic number to speed up deallocation
- Embedding free list
 - Build free list inside the free space itself
 - Nodes with size and next pointer
- Growing heap
 - Just give up and return NULL
 - Or call sbrk system call to OS to grow heap

11.3 Basic Strategies

- Best fit
 - Return smallest chunk that is equal or larger than the requested size
 - Requires linear search
- Worst fit
 - Find largest chunk, split it, return requested size
 - Requires linear search
- First fit
 - Returns first block big enough
 - Faster because no exhaustive search
- Next fit
 - Returns first block big enough starting from previous location
 - Spreads searches through free space more uniformly

11.4 Other Approaches

- Segregated Lists
 - Keep separated list to manage all objects of that size
 - Hard to determine much memory to dedicate to that list
- Slab allocator by Jeff Bonwick
 - Object caches for kernel objects
 - Each object cache are segregated free lists
 - Requests slabs of memory from general allocator, when running low
- Binary buddy Allocation
 - Big space of 2^N
 - Suffers from internal fragmentation but can recursively coalesce

12 L5: Garbage Collection and Defragmentation

12.1 Garbage Collection

- Allocated resources are freed through explicit/implicit action by client
 - `close(2)`, `free(3)`, delete operator, returning from a C/C++ subroutine, `exit(2)`
- If shared by multiple concurrent clients
 - Free only if reference count is zero (don't free if others are still using it, just decrement the reference count)
- Garbage Collection
 - Analyzes allocated resources to determine which are still in use
 - Data structures assoc with resource references are designed to be easily enumerated to enable the scan for accessible resources
 - Comes at a performance cost

12.2 Defragmentation

- Shards of free memory are not useful
 - Coalescing is only useful if adjacent memory free at same time
- Defragmentation
 - Changes which resources are still allocated
- Flash management
 - NAND Flash is a pseudo-Write-Once-Read-Many medium
 - Identify large (64MB) block with many 4KB blocks not in use
 - Move all in use blocks and update resource allocation map
 - Erase large block and add 4KB blocks to free list
- Disk Space Allocation
 - Choose region to create contiguous free space
 - For each file in that region, move it elsewhere
 - Coalesce all that free memory
 - Move set of files into that region
 - Repeat until all files and free space is contiguous
- Internal fragmentation is like rust, it never sleeps
 - Defragmentation used to be run periodically, now is run continuously
- Conclusions
 - If using garbage collection, must make all resources discoverable, how to trigger scans, prevent race conditions with application
 - Must not disrupt running applications when using defragmentation

13 L6: Arpaci-Dusseau Chapter 18: Paging: Introduction

13.1 A Simple Example And Overview

- Paging
 - Divide process address space into fixed-sized units
 - View memory as fixed-sized page frames
- Free list
 - OS may hold list of free pages
- Page table is a per process data structure
 - Stores address translations for virtual pages so we know where it is in physical memory
- Virtual address [VPN, OFFSET]
 - Virtual page number (VPN) indexes page table to find physical frame/page number (PFN/PPN)
 - Translate VPN to PPN then load from memory
 - Offset determines which byte within page

13.2 Where Are Page Tables Stored?

- Page table entry
 - Holds physical translation
 - If roughly 4 bytes per PTE, page tables would be big
 - Problem bc page table per process
- Stored somewhere in memory

13.3 Whats Actually In The Page Table?

- Linear Page Table
 - Index array by VPN to look up PTE and to find physical frame number (PFN)
 - Valid bit indicates if memory is valid (traps if invalid)
 - Proction bit indicates whether page can be read/written/executed (trap if bad access)
 - Present bit indicates whether page is in memory or disk (if it has been swapped out)
 - Dirty bit indicates if page has been modified since brought into memory
 - Reference/aceess bit indicates if page has been accessed (to determine which pages are popular; used for page replacement)

13.4 Paging: Also Too Slow

- Must translate virtual address
 - $VPN = (\text{Virtual address} \& VPN_{MASK}) \gg \text{SHIFT}$
 - $PTEaddr = \text{Page table base address} + VPN * \text{sizeof}(PTE)$
 - $Offset = \text{Virtual address} \& OFFSET_{MASK}$
 - $PhysAddr = (PFN \ll \text{SHIFT}) \mid Offset$

14 L6: Arpaci-Dusseau Chapter 19: Paging: Faster Translations (TLBs)

14.1 TLB Basic Algorithm

- TLB
 - Bc chopped address space into many fixed-sized units, paging requires a lot of memory to map addresses
 - This mapping memory is also stored in physical memory, which would require an additional memory lookup to read
 - Instead, use a TLB, which is an address translation cache, to hold popular virtual-to-physical translations
- TLB Hit/miss
 - If virtual page number (VPN) from virtual address (VA) is inside the TLB (translation lookaside buffer), then have TLB hit and may extract the page frame number (PFN)

- If VPN from VA is not inside TLB, then have TLB miss and must access page table (in memory) to find translation, update TLB, then restart lookup into TLB

14.2 Example: Accessing An Array

- Start with a miss, then multiple hits
 - Rely on spatial locality for first pass
 - Rely on temporal locality for second pass
- Caching is fundamental
 - Temporal and spatial locality are necessary
 - Can't make caches large because physics says large cache is slow

14.3 Who Handles The TLB Miss?

- Hardware
 - Use page table base register to walk page table and find PTE
- Software
 - Hardware raises exception, pauses instructions, privilege raises to kernel mode, jumps to trap handler
- Infinite TLB misses
 - If is a problem, keep TLB miss handlers in physical memory (unmapped) so it will always be a hit
- RISC vs CISC (Aside)
 - Complex has more and more powerful instructions
 - Reduced has fewer and simpler primitives

14.4 TLB Contents: Whats In There?

- Fully associative means a given translation can be anywhere in the TLB
- VPN — PFN — other bits
 - Other bits include valid bit, protection bits (regarding w/r/x), address space identifier, dirty bit, etc

14.5 TLB Issue: Context Switches

- Fully associative means a given translation can be anywhere in the TLB
- VPN — PFN — other bits
 - Other bits include valid bit, protection bits (regarding w/r/x), address space identifier, dirty bit, etc
 - Could flush TLB on context switch, or could use address space identifier

14.6 Issue: Replacement Policy

- LRU Replacement Policy
 - Least recently used, but usually can't actually do this, so vaguely do LRU

15 L6: Arpaci-Dusseau Chapter 21: Beyond Physical Memory: Mechanisms

15.1 Swap Space

- Swap Space
 - Use hard disk drive as storage
 - Reserved space on disk for moving pages back and forth

15.2 The Present Bit

- Extract VPN from VA, check for TLB hit and produce PA if possible
 - Otherwise, receive TLB miss and go to memory through page table base register to find PTE
- Present bit
 - Set to one if page is in physical memory
 - Otherwise, is not in physical memory and is a page fault
 - OS invoked to service page fault, so page-fault handler runs

15.3 The Page Fault

- OS page-fault handler
 - Hardware does not do it because hardware does not know enough about swap space, I/O, etc
 - OS looks in PTE to find address and request it from disk
 - Process is blocked during this, so run another process

15.4 What If Memory Is Full?

- Page-replacement Policy
 - Page in from swap space; Page out from memory
 - Replace if memory is full
 - 10k-100k times slower if poor page-replacement policy

15.5 Page Fault Control Flow

- If TLB miss
 - If invalid, OS trap handle terminates process
 - If not present, run page fault handler
 - Find physical frame for soon-to-be-faulted-in page
 - Run replacement alg if necessary
 - I/O request page from swap space
 - Retry for TLB miss, then retry for TLB hit
 - If present and valid, grab PFN from PTE and retry

15.6 When Replacements Really Occur

- Swap daemon
 - If fewer pages than the low watermark, then background thread evicts pages
 - Continues evicting pages until the high watermark
 - Then goes back to sleep and waits
- Clustering
 - Clustering/grouping these pages to swap partition increases efficiency because it reduces disk seek and rotational overheads
- Background work
 - Do work in background (buffered disk writes, etc) because it is more efficient and makes better use of idle time

16 L6: Arpaci-Dusseau Chapter 22: Beyond Physical Memory: Policies

16.1 Cache Management

- Minimize cache misses because a single miss will make it very slow
 - Average memory access time (AMAT) = $T_M + (P_{miss} * T_D)$

16.2 The Optimal Replacement Policy

- Farthest in the future
 - Is optimal
 - Use this as a reference point, something to compare our algorithms against
- Types of misses
 - Cold-start miss is compulsory because cache is empty
 - Capacity miss is because cache ran out of space
 - Conflict miss is because of hardware limits on where items can be placed in a hardware cache (not a problem for OS page cache)

16.3 Replacement Policies

- FIFO
 - Performs quite terribly, but is simple to implement
 - Belady's Anomaly: FIFO performs even worse on larger cache than on smaller cache
- Random
 - Can work
- Least-Frequently-Used (LFU)/Least-Recently-Used (LRU)
 - Rely on locality and do what their names say
- Most-Frequently-Used (MFU)/Most-Recently-Used (MRU)
 - Exist and do not work well
- Workload examples
 - FIFO doesn't do well, random can do well, LRU does fairly well

16.4 Implementing LRU

- True LRU is expensive
 - Finding truly least-recently-used page is prohibitively time-consuming
- Approximate LRU using Clock algorithm
 - Whenever page is referenced, use bit is set
 - Clock hand points to some page, if bit is set, unsets it and checks next
 - If bit is unset, replaces it

16.5 Considering Dirty Pages

- If page is dirty (set dirty bit), then must be written back to disk if we want to evict it
 - Prefer to evict clean pages

16.6 Other VM Policies

- Demand Paging
 - Bring page into memory only 'on demand'
 - Opposite of prefetching memory
- Clustering/Grouping of writes
 - Write many things at same time because of how disk drive works

16.7 Thrashing

- If memory is just oversubscribed
 - Then will constantly page and thrash
- Admission control
 - Decide to not run some processes, so that we may do well on the remaining processes
- Out-of-memory killer
 - Will choose a memory-intensive process and kill it

17 L6: Working Sets

17.1 LRU is not enough

- Global LRU
 - Most-recently used page is from current process and will not run for a while
 - Least-recently used page is from old process about to run

17.2 The concept of a Working Set

- Is the set of pages for a given process
 - Increasing the number of pages makes little difference in performance, but decreasing makes a difference
- Different computations require different sizes, getting the number correct will minimize page faults and maximize throughput

17.3 Implementing Working Set replacement

- More information recorded about pages
 - Associated with owning process
 - Accumulated CPU time
 - Last referenced time
 - Target age parameter
- Age decisions are made on the basis of accumulated CPU time
 - Page ages if owner runs without them
 - Pages younger than a target age are preferably not replaced
 - Give pages older than target age away

17.4 Dynamic Equilibrium to the rescue

- Page stealing algorithm
 - Every process is continuously losing and stealing pages
 - Processes that reference more pages more often will accumulate larger working sets while others will find their set reduced
 - Working sets adjust automatically

18 L7: Arpaci-Dusseau Chapter 25: A Dialogue on Concurrency

18.1 Dialogue

- Multi-threaded applications
 - Threads access memory; we don't want multiple threads to access memory at same time
 - OS supports primitives such as locks and condition variables

19 L7: Arpaci-Dusseau Chapter 26: Concurrency: An Introduction

19.1 Introduction

- Context switch
 - Save state (program counter, registers) to thread control block
 - Address space stays the same, so page table does not need to be switched
- Multiple stacks in address space if multiple threads
- Thread-local storage
 - Stack of that thread

19.2 Why Use Threads?

- Used threads to exploit parallelism
 - If single processor, then not relevant
 - Otherwise, parallelize and used thread per CPU
- Use threads to do something when blocked program
 - Instead of waiting for IO, just switch to another thread and do things
- Choose process for logically separate tasks with little sharing of data structures

19.3 An Example: Thread Creation

- Use pthreads
- Will run in different order according to scheduler
- May not be deterministic

19.4 The Heart Of The Problem: Uncontrolled Scheduling

- Race condition
 - Execution depends on timing execution of code (indeterminate)
- Critical section
 - Multiple threads executing code resulting in race condition
- Mutual exclusion
 - If one thread executing inside critical section, others will be prevented

19.5 The Wish For Atomicity

- Atomic (all or nothing)
 - Don't just have atomic instructions for all because too many instructions
- Synchronization primitives
 - General set of instructions to control multi-threaded programs

20 L7: Arpaci-Dusseau Chapter 27: Interlude: Thread API

20.1 Threads

- Use `pthread_create` to create new thread
- Use `pthread_join` to wait for thread to complete
- Use `pthread_mutex_lock` and `pthread_mutex_unlock` to provide mutual exclusion to critical sections via locks
 - Need to properly initialize and check that lock/unlock actually succeed
- Condition variables
 - Enables thread to wait until particular condition occurs
 - Needs lock and condition
 - Sleeps until other thread signals
- Spinlock
 - Wait in loop until lock available, consuming CPU cycles

21 L7: User-Mode Thread Implementation

21.1 Introduction

- Threads are independent schedulable unit of execution
 - Runs within address space of process
 - Has access to system resources from process
 - Has own registers and stack

21.2 User/Kernel

- User-level thread done without OS
 - Allocates memory, dispatches thread, sleeps, exits, free memory
 - If system call blocks, entire process blocks
 - Cannot exploit multi-processors
- Kernel implemented threads
 - Exploits multi-processors and switches between threads when one blocks

21.3 User/Kernel

- Non-preemptive scheduling
 - User-mode threads are more efficient than kernel for contex-switches
- Preemptive scheduling
 - Allowing OS to schedule is better than setting alarms and signals

22 L7: Inter-Process Communication

22.1 Introduction

- coordination of operations with other processes
 - synchronization (e.g. mutexes and condition variables)
 - the exchange of signals (e.g. `kill(2)`)
 - control operations (e.g. `fork(2)`, `wait(2)`, `ptrace(2)`)
- the exchange of data between processes:
 - uni-directional/bi-directional

22.2 Simple Uni-Directional Byte Streams

- Pipes
 - Opened by parent and inherited from child
 - Each program in pipeline is unaware of what others do, byte streams are unstructured, etc
 - If reader exhausts data in pipe, reader does not get EOF (is blocked instead)
 - Flow control: Available buffering capacity of pipe may be limited, so writer may be blocked for reader to catch up
 - Writing to pipe without open read fd is illegal (gets signal exception)
 - When both read/write fd are closed, pipe file is deleted
- Only data privacy mechanisms are on initial/output file
 - Generally no auth/encryption while passing

22.3 Named Pipes and Mailboxes

- Named-pipe fifo
 - Persistent pipe whos reader/writers can open by name (rather than inheriting)
 - Writes may be interspersed
 - Readers/writers can't authenticate identity
- Mailboxes
 - Data is not bytestream, each write is stored as message
 - Each write has authenticated ID
 - Unprocessed msgs remain in mailbox

22.4 General Network Connections

- Higher level communication/service models
 - Remote procedure calls - distributed request/response APIs
 - RESTful service models - layered on HTTP GETS/PUTS
 - Publish/Subscribe services - content based info flow
- Complexity
 - Interoperability with software running different OS and ISA
 - Security issues, changing addresses, failing connections

22.5 Shared Memory

- High performance for Inter-Process Communication
 - Efficiency wrt low cost per byte
 - Throughput wrt bytes per second
 - Latency wrt minimum delay
- Ultra high performance
 - Shared memory by creating a file for communication
 - Process maps file into virtual address space
 - Is available immediately upon writing
 - Very fast but can only be used on same memory bus
 - Has no authentication and a single bug can kill both

22.6 Network Connections and Out-of-Band Signals

- Preempting queued operations
 - Have a reserved out-of-band channel so signal can preempt others if urgent
 - Adds overhead but allows important messages to skip FIFO line (network connection is FIFO)

23 L7: Named pipes, Send, Recv, Mmap

23.1 Named Pipes

- Named pipes exist as device special file
- Can be accessed by processes of different ancestries
- When I/O done, pipe remains
- Normally, if FIFO opened for reading, process will block until another process opens it for writing
- If write to pipe without reader, will get SIGPIPE

23.2 Send

- Send a message on a socket

23.3 Recv

- Receive a message from a socket

23.4 Mmap

- Map or unmap files or devices into memory
- Creates a new mapping in the virtual address space of the calling process

24 L8: Arpaci-Dusseau Chapter 28: Locks

24.1 Locks: The Basic Idea

- Use of lock
 - Put around critical sections so that it is performed atomically
- Lock variable
 - If no other thread holds lock, thread acquires lock and enters critical section
 - If another thread holds lock, then will not return

24.2 Pthread Locks

- Mutex is the POSIX library lock
 - Provides mutual exclusion (exclude other threads from entering until first thread has completed)
- Use multiple locks (as opposed to one big lock for any critical section)
 - Fine-grained vs coarse-grained approach

24.3 Evaluating Locks

- Mutual exclusion, fairness, performance
 - Needs to prevent multiple threads from entering critical section
 - Needs to not let contending threads starve
 - Need time overheads to not be high

24.4 Controlling Interrupts

- Disable interrupts during critical section to provide mutual exclusion
 - For single-processor system, makes code atomic
- Cons
 - Requires user to call privileged operation
 - Greedy user could lock for entire process
 - Buggy user could break computer
 - Does not work for multiprocessors because multiple threads can still enter critical section

- Interrupts may be lost
- OS is allowed to use this as mutual-exclusion primitive for updating data structures

24.5 A Failed Attempt: Just Using Loads/Stores

- Flag
 - Doesn't work because the checking/setting of flag is not atomic
 - Also spin-waiting (persistently checking value of flag) is incredibly inefficient

24.6 Building Working Spin Locks with Test-And-Set

- Test-and-set instruction (atomic exchange)
 - Puts new value into old value; returns old value (atomically)
 - Is sufficient to build a spinlock
- Spinlock
 - To work on a single processor, requires preemptive scheduler (otherwise, thread would never relinquish CPU)
 - Is a correct lock
 - No fairness guarantees
 - Terrible performance if single processor
 - If N threads contending, N-1 time slices may be wasted while spinning on single processor
 - Okay performance if multiple processors

24.7 Compare-And-Swap

- Compare-and-swap (compare-and-exchange)
 - Test if value of ptr is equal to value at expected, if so, update with new value, otherwise, do nothing
 - Can build spinlock with this

24.8 Load-Linked and Store-Conditional

- Load-linked
 - Fetch value from memory and put in register
- Store-conditional
 - If success, updates and returns 1; otherwise, no update and returns 0
 - Only one thread is able to acquire lock if using these (because store-conditional will fail)

24.9 Fetch-And-Add

- Fetch-and-add
 - Increment value and return old value
- Ticket lock
 - If thread wants lock, do fetch-and-add and wait
 - Global lock-;turn determines who's turn
 - All threads make progress

24.10 A Simple Approach: Just Yield, Baby

- Yield
 - System call yield to allow processes to deschedule self
 - Better than spinlock, but still costly
 - Does not address starvation issue

24.11 Using Queues: Sleeping Instead Of Spinning

- Sleep and wake
 - Test-and-set with explicit queue of lock waiters
 - Avoids starvation
 - May sleep forever in wakeup/waiting race if release of lock occurs after park()
 - So have setpark() to indicate a thread is about to park; if interrupted and another thread unparks, thread will return rather than sleep
- Solaris uses park/unpark
- Linux uses futex

24.12 Two-Phase Locks

- Spins during first cycle, then on second cycle will sleep
- Hybrid approach is effective

25 L6: Arpaci-Dusseau Chapter 30.1: Condition Variables

25.1 Definition and Routines

- Condition variable
 - Explicit queue for threads if condition is not met
 - When condition is correct, wakes

26 L6: Inter-Process Communication

26.1 Introduction

–